

Part 8: File Systems

- File System Structure
- Logical File System
 - Files: File Attributes, File Types, File Structure, File Access Methods, File Operations
 - Directories: Directory Structure, Directory Organisation, Directory Operations
 - File Protection in UNIX
 - In-memory File System Data Structures
- File Organization Module
 - Allocation Methods: Contiguous Allocation, Linked Allocation (& FAT), Indexed Allocation (& inode)
 - Disk-Space Management: Block Size, Keeping Track of Free Blocks

File System Structure

- A file system is generally composed of many different levels. For example:
 - *Logical File System*: manages directory structure, responsible for file creation, access, deletion, protection and security.
 - *File-Organisation Module*: allocates storage space for files, translates logical block addresses to physical block addresses, and manages free disk space.
 - *Basic File System*: manages buffers and caches issues generic commands to the appropriate device driver to read and write physical blocks on the disk.
 - *I/O Control*: consists of device drivers and interrupt handlers to transfer information between memory and the disk system.

File Attributes

- A file may have the following attributes
 - *Name* - only information kept in human-readable form.
 - *Type* - needed for systems that support different types.
 - *Location* - pointer to file location on device.
 - *Size* - current file size
 - *Protection* - controls who can do reading, writing, executing.
 - *Time, Date, and User Identification* - data for protection, security, and usage monitoring.
- These information about files are kept in the directory structure, which is maintained on the disk.

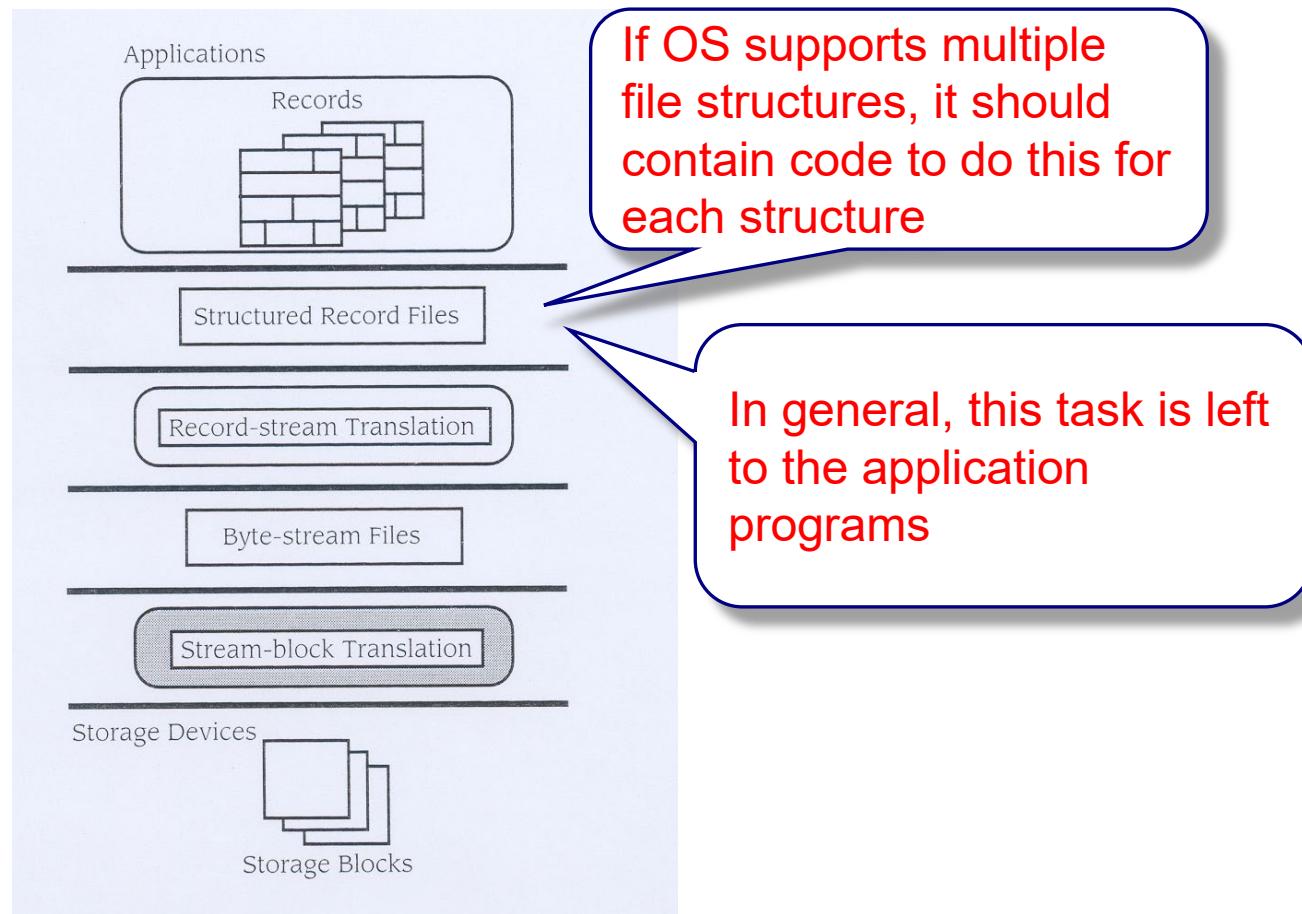
File Types

- A file has a contiguous logical address space which can store many different types of information:

File Type	Usual extension	Function
Executable	exe, com, bin or none	ready-to-run machine-language program
Object	obj, o	compiled, machine language, not linked
Source code	c, p, pas, 177, asm, a	source code in various languages
Batch	bat, sh	commands to the command interpreter
Text	txt, doc	textual data documents
Word processor	wp, tex, rrf, etc.	various word-processor formats
Library	lib, a	libraries of routines
Print or view	ps, dvi, gif	ASCII or binary file
Archive	arc, zip, tar	related files grouped into one file, sometimes compressed.

File Structure

- A file consists of a collection of records. These records can be organised or structured to facilitate file access.



File Structure (Cont.)

- Unstructured: Sequence of Bytes.
 - A file is a stream of bytes. Each byte is individually addressable from the beginning of the file.
 - Used by UNIX and MSDOS (and **assumed in the following discussions**)

File Access Methods

- Sequential Access: Information in a file is processed in order from the beginning of the file, one **byte** after the other.
- Direct Access: **Bytes** of a file can be read in any order (by referencing **byte** number).

File Operations

Commands	Explanation
Create	allocate disk space; create directory entry with file attributes.
Delete	delete the corresponding directory entry; deallocate disk space.
Open	search the directory structure for file entry; move the content to memory (put in the <i>open file table</i>). <u>Information Associated with an open file:</u> - <i>Current Position Pointer</i> - <i>File Open Count</i> - <i>Disk Location</i>
Close	move the content of directory entry in memory to directory structure on disk.
Write	search <i>open file table</i> to find the location of file; write data to the position pointed by <i>Current Position Pointer</i> .
Read	search <i>open file table</i> to find the location of file; read data from the position pointed by <i>Current Position Pointer</i> .

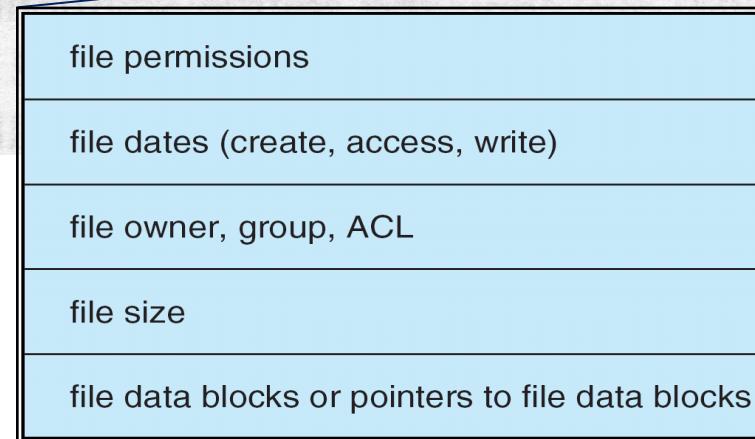
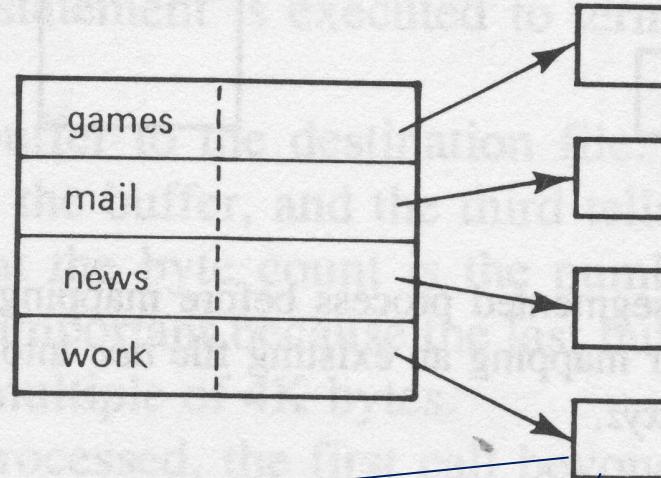
Directory Structure

- A directory typically contains a number of entries, one per file. Both the directory and the files reside on disk.
- A directory can be structured in two ways:
 - (a) each entry contains a file name and other attributes; or
 - (b) each entry contains a file name and a pointer to another data structure where file attributes can be found.

Directory Structure (Cont.)

games	attributes
mail	attributes
news	attributes
work	attributes

(a)

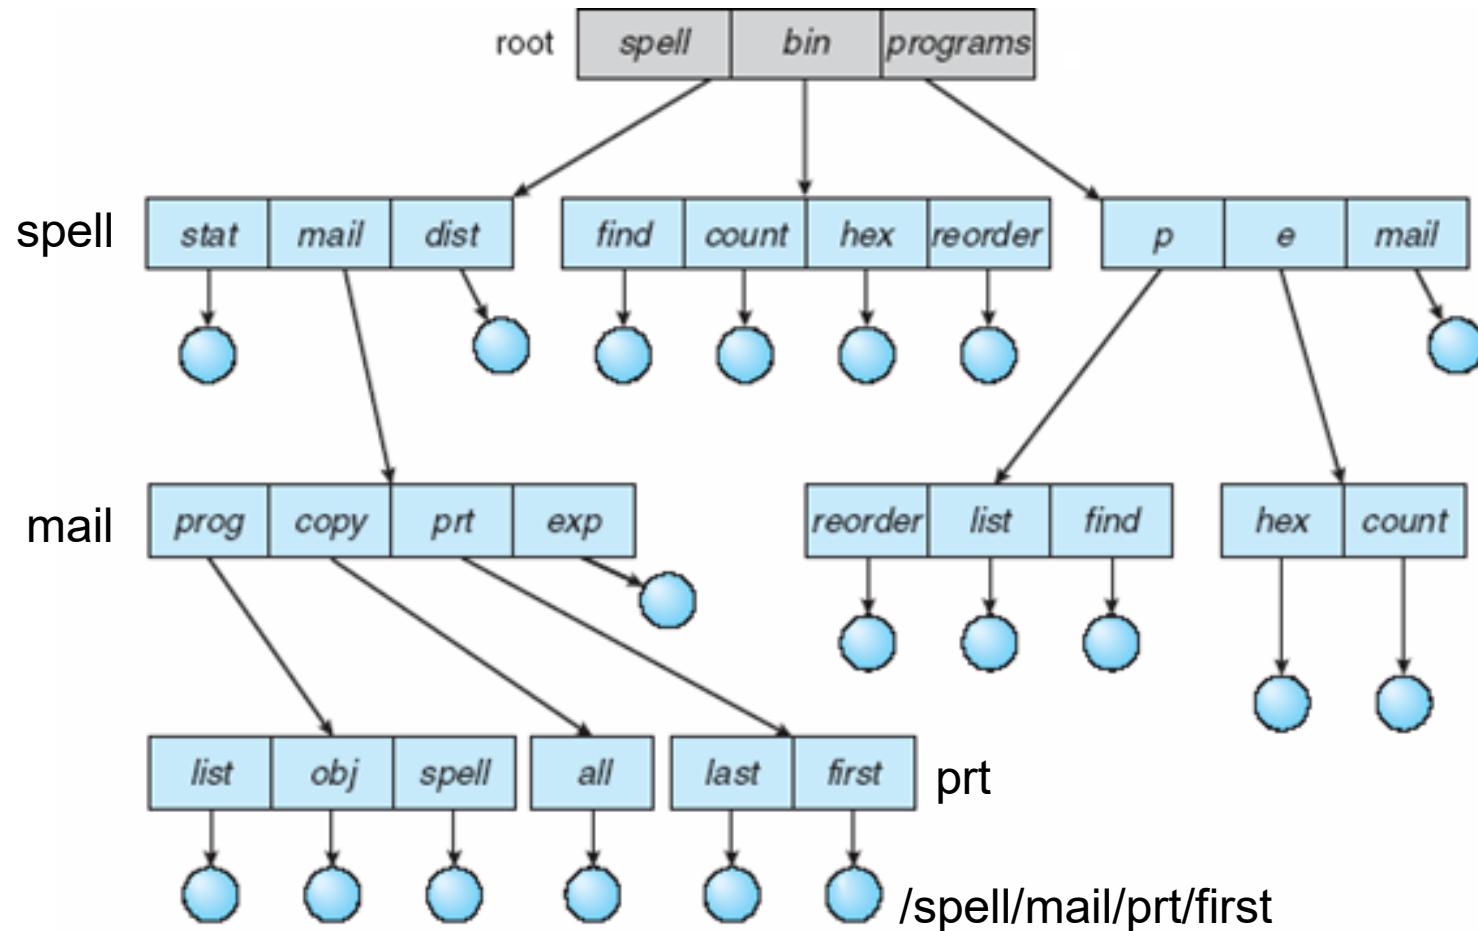


**File Control Block
(FCB)**

Directory Structure (Cont.)

- A directory can be implemented using:
 - *Linear List*: Directory entries are arranged as linear list
 - simple to implement
 - time-consuming to execute: require linear search to find a particular entry
 - *Hash Table*: Directory entries are arranged as linear list with a hash table to facilitate file look-up.
 - decreases directory search time
 - collisions: situations where two file names hash to the same location (solution: each hash table entry contains a list of directory entries)

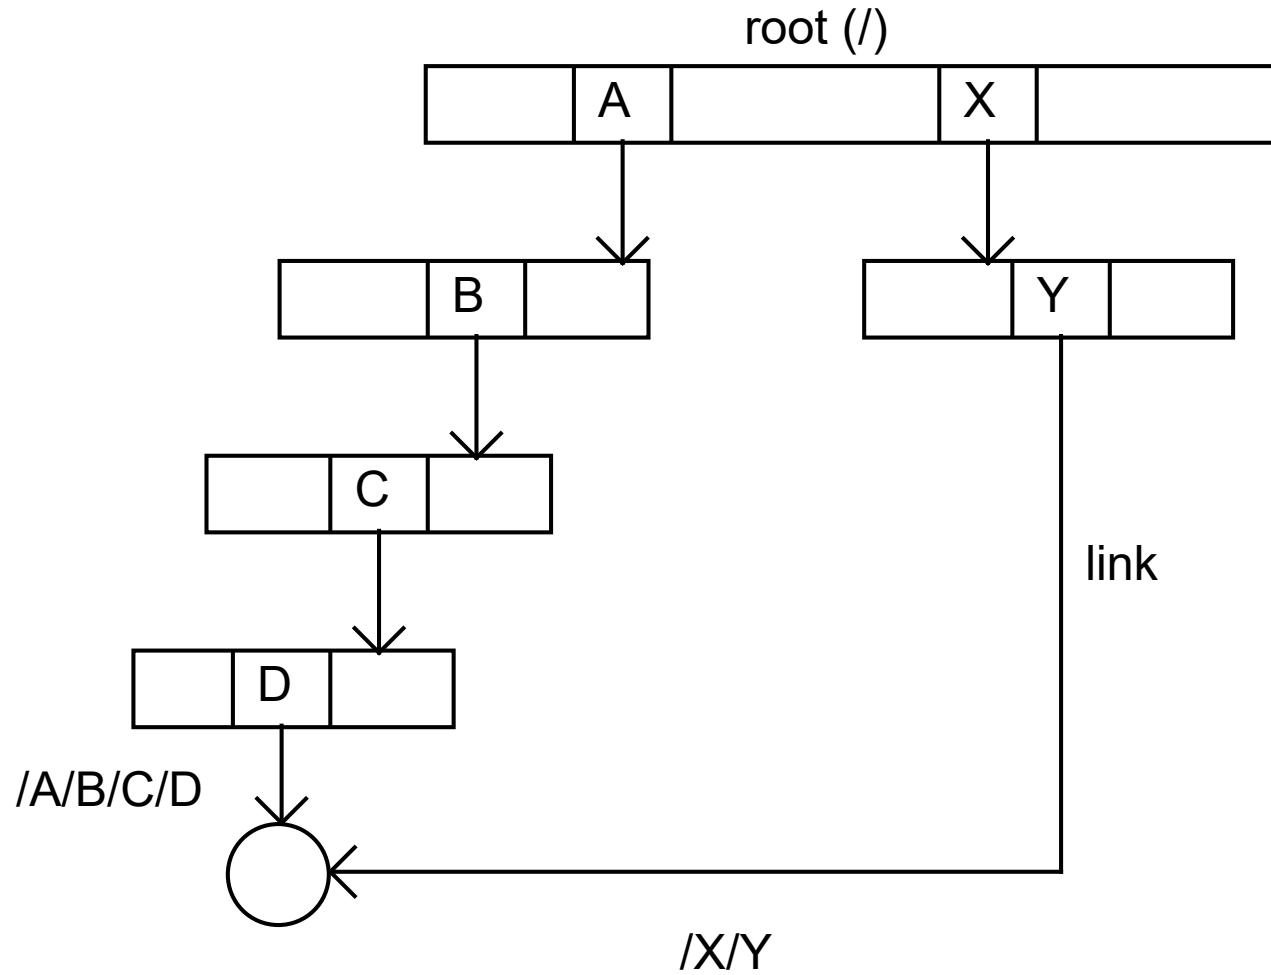
Directory Organization



Directory Organization (Cont.)

- Path Name
 - *Absolute Path Name*: begins at the root and follows a path down to the specific file, e.g., /spell/mail/prt/first
 - *Relative Path Name*: Defines a path from the current directory, e.g.,
 - current directory is: /spell/mail
 - relative path name for the above file is: prt/first
- Characteristics
 - *Efficient Searching*: File can be easily located according to the path name.
 - *Naming*: Files can have the same name under different directories.
 - *Grouping*: Files can be grouped logically according to their properties.

Acyclic-Graph Directory



Acyclic-Graph Directory (Cont.)

- Support for File Sharing: Two methods

1. *Symbolic Link*

- create a directory entry *link*, which contains absolute or relative path name of a file
- resolve the link by using the path name to locate the real file
- slower access than with hard link

2. *Hard Link*

- duplicate all information about a file in multiple directories

Acyclic-Graph Directory (Cont.)

- Problems with File Sharing:

- In traversing the file system, the shared files may be visited more than once

Solution: Ignore the link entry when traversing

- Deleting a shared file may leave dangling pointers to the now-nonexistent file

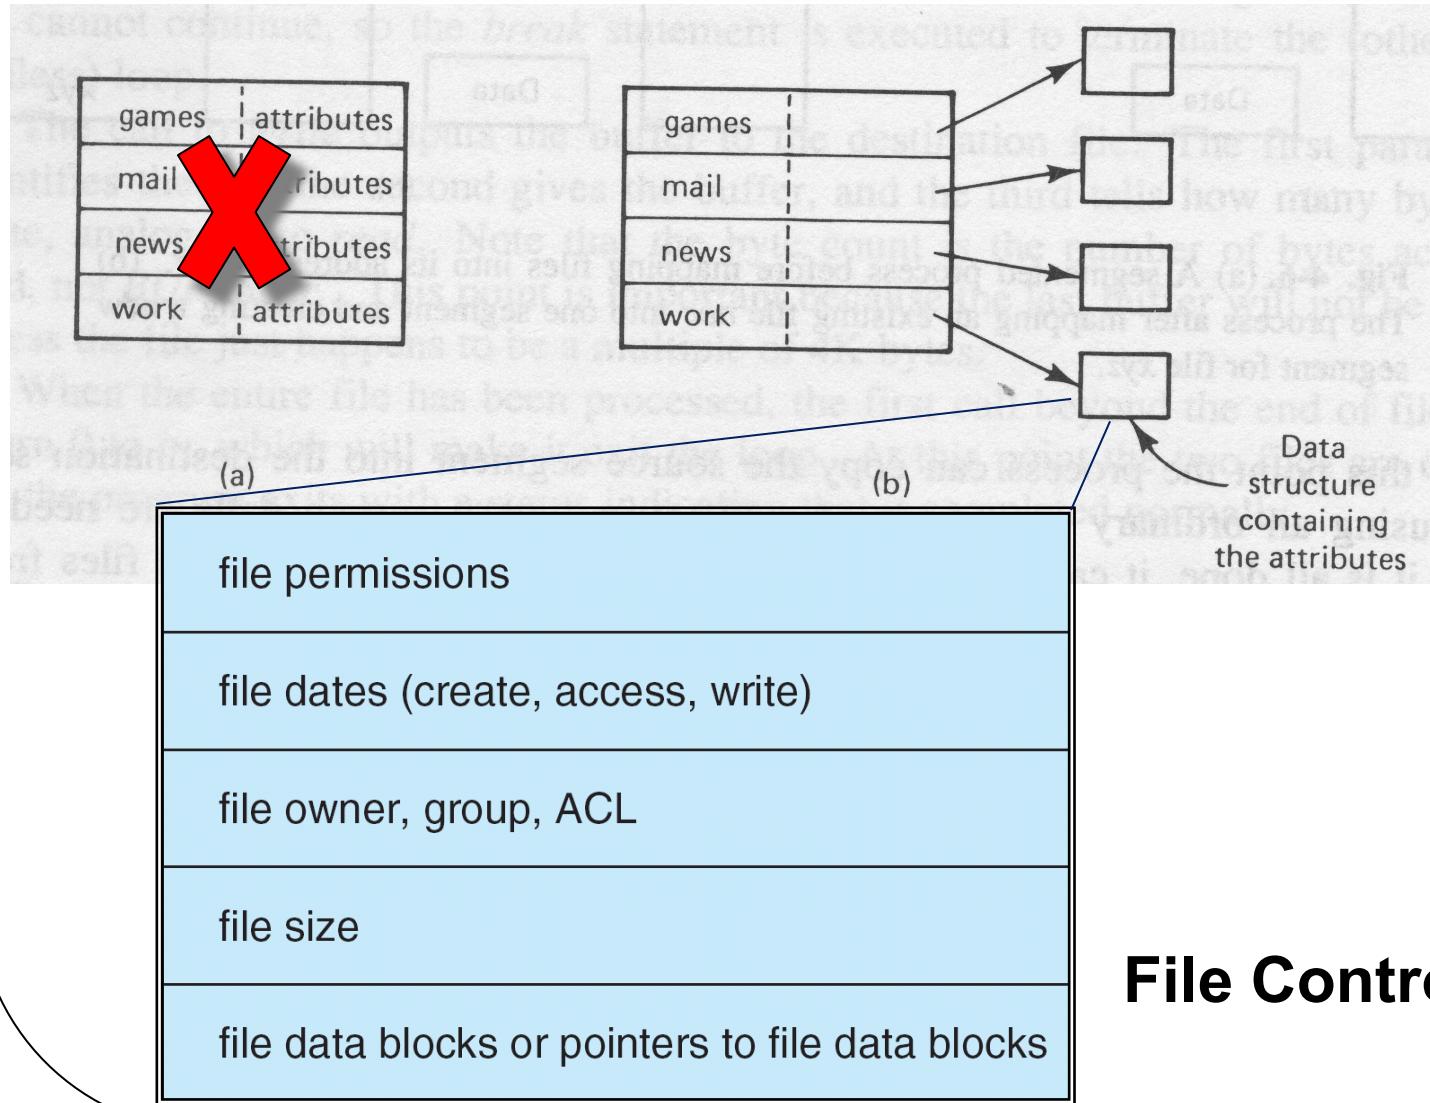
Solution:

- search for dangling links and remove them; or
 - leave the dangling links and delete them only when they are used again; or
 - preserve the file until all references to it are deleted.

Directory Operations

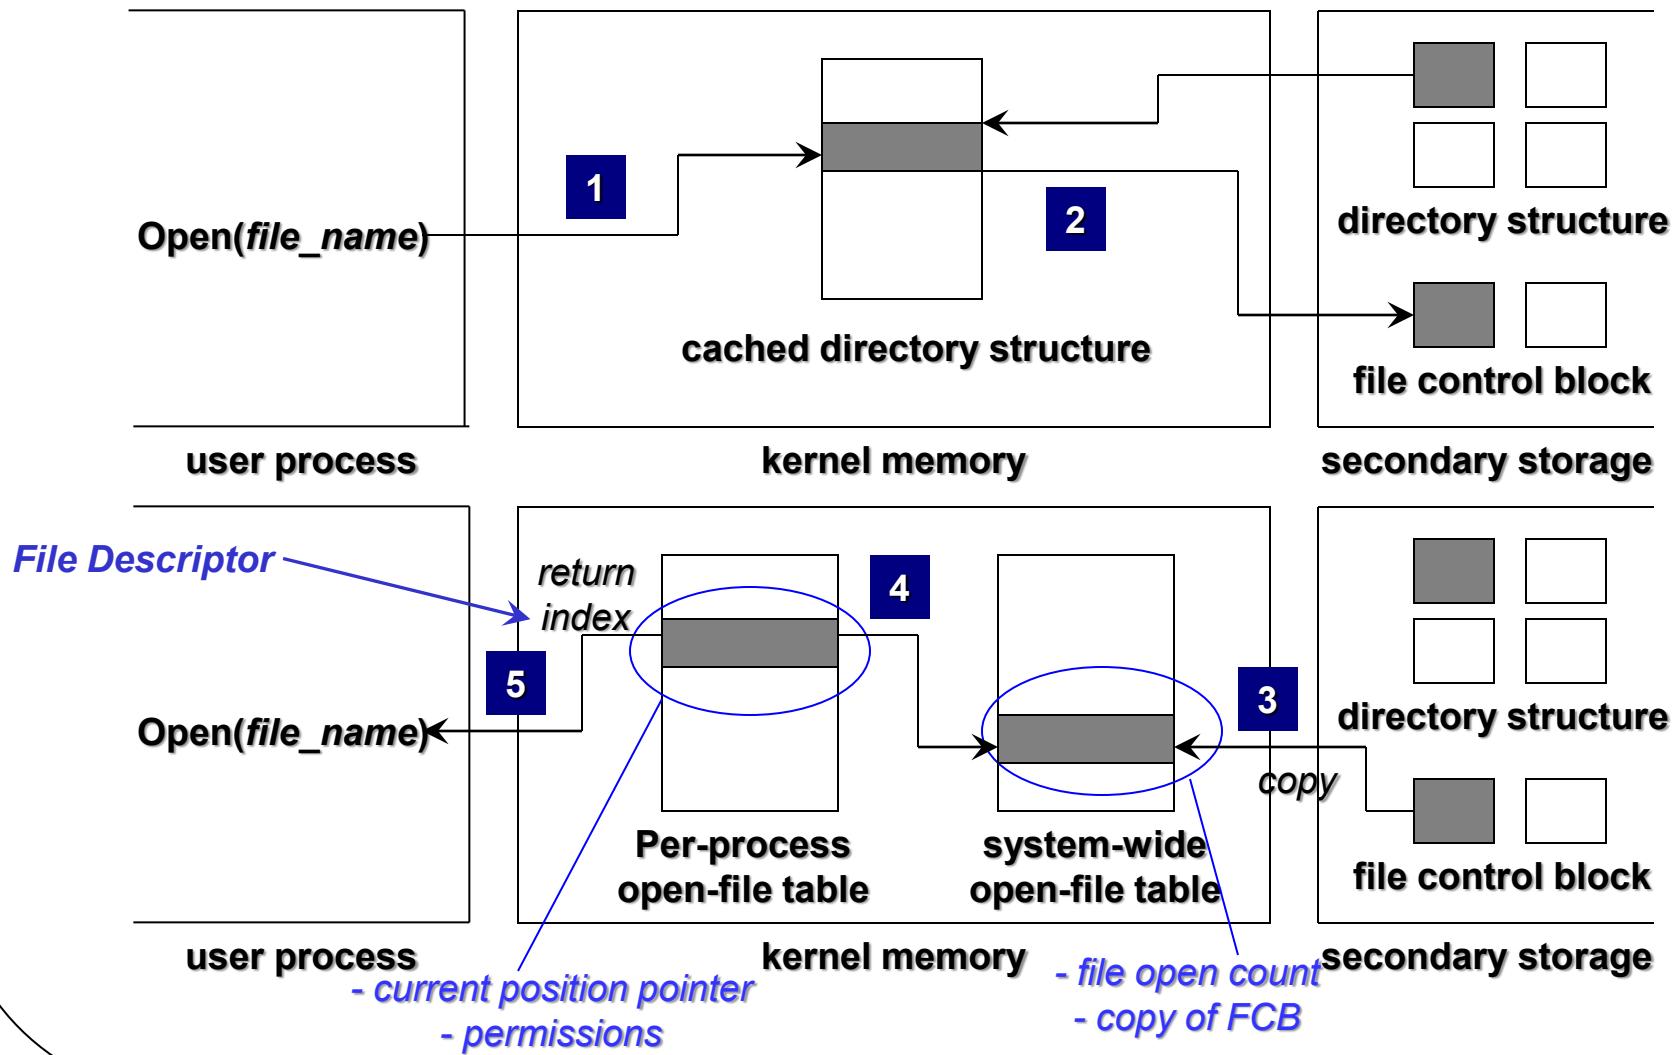
<i>Commands</i>	<i>Explanation</i>
Create	create a directory. In UNIX, two entries "." and ".." are automatically added when a directory is created. "." refers to the <i>current directory</i> ; and ".." refers to its <i>parent</i> .
Delete	delete a directory. Only empty directory can be deleted (directory containing only "." and ".." is considered empty).
List	list all files (directories) and their contents of the directory entry in a directory
Search	search directory structure to find the entry for a particular file.
Traverse	access every directory and every file within a directory structure.

In-memory File System Data Structure

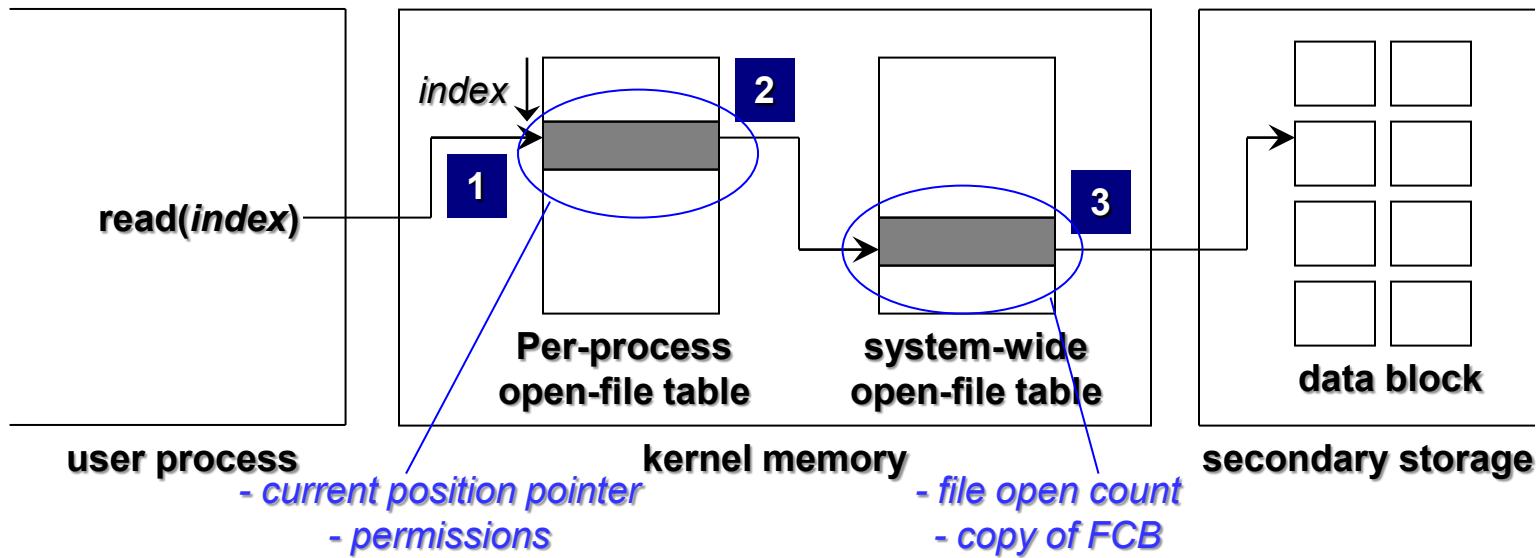


File Control Block

In-memory File System Data Structure



In-memory File System Data Structure



- *Current position pointer & amount of data to read ⇒ logical file blocks to be assessed*
- *Information in FCB & file allocation method ⇒ where these blocks are on hard-disk*

File Protection in UNIX

- Model of Access: *read, write, execute*
- Three Classes of Users:

		RWX
– owner access:	7	⇒ 1 1 1
– group access:	5	⇒ 1 0 1
– public access:	1	⇒ 0 0 1

File type	User access	Group access	Public access	Links	Userid	Size	Date	Time	File name
↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
-	r w x	r - x	- - x	1	smith	58	Mar 3	3:04	game1

File Protection in UNIX

- Meaning of Permissions for a Directory:
 - To access a directory, the execute permission is essential. No execute permission,
 - can't execute any command on the directory
 - have no access to any file contained in the file hierarchy rooted at that directory
 - No read permission -> can't list the directory
 - No write permission -> can't create or delete files in the directory

File type	User access	Group access	Public access	Links	Userid	Size	Date	Time	File name
↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
d	r w x	r - x	---	0	smith	5	Feb 2	12:01	games

File System Structure

- A file system is generally composed of many different levels. For example:
 - *Logical File System*: manages directory structure, responsible for file creation, access, deletion, protection and security.
 - *File-Organisation Module*: allocates storage space for files, translates logical block addresses to physical block addresses, and manages free disk space.
 - *Basic File System*: manages buffers and caches issues generic commands to the appropriate device driver to read and write physical blocks on the disk.
 - *I/O Control*: consists of device drivers and interrupt handlers to transfer information between memory and the disk system.

Allocation Methods

- Contiguous Allocation
- Linked Allocation
- Indexed Allocation

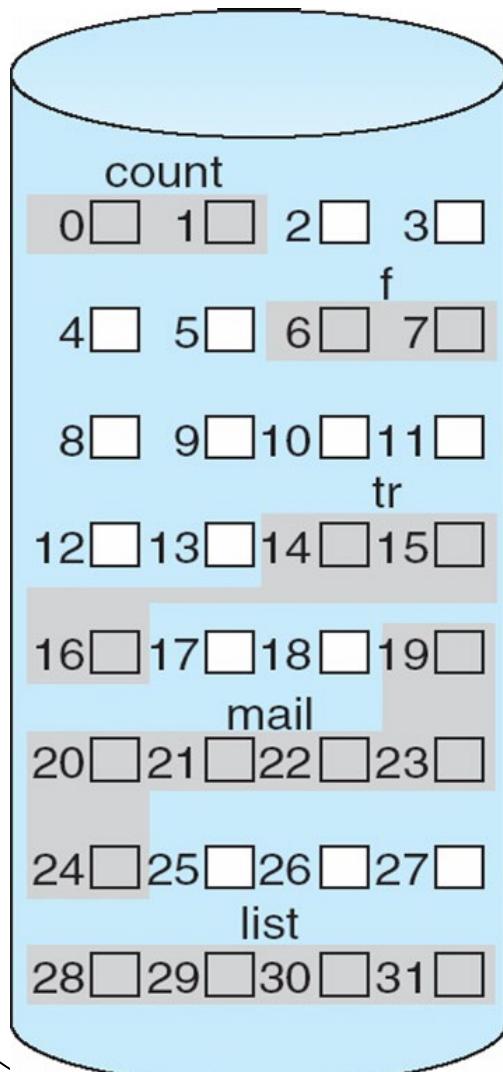
Contiguous Allocation

- Each file occupies a set of contiguous blocks on the disk.
 - Simple: only starting location (block #) and length (number of block) are required.
 - Support random access
- Problems:
 - Waste of space (similar to dynamic storage allocation of main memory)
 - Finding hole big enough using First-fit (faster) or Best-fit may result in external fragmentation
 - File space constricted by size of hole, so may later have to move to a bigger hole
 - If instead needed file space is overestimated \Rightarrow internal fragmentation

Contiguous Allocation (Cont.)

- Logical to Physical Address Mapping:
 - Suppose block size is 512 bytes
 - Logical address $\Rightarrow Q \times 512 + R$
 - Block to be accessed = $Q + \text{starting address}$
 - Displacement into block = R

Contiguous Allocation (Cont.)



directory

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2

Logical Address = 2333

$$= Q \times 512 + R$$

$$= 4 \times 512 + 285$$

Physical Address = 285th byte of block
23

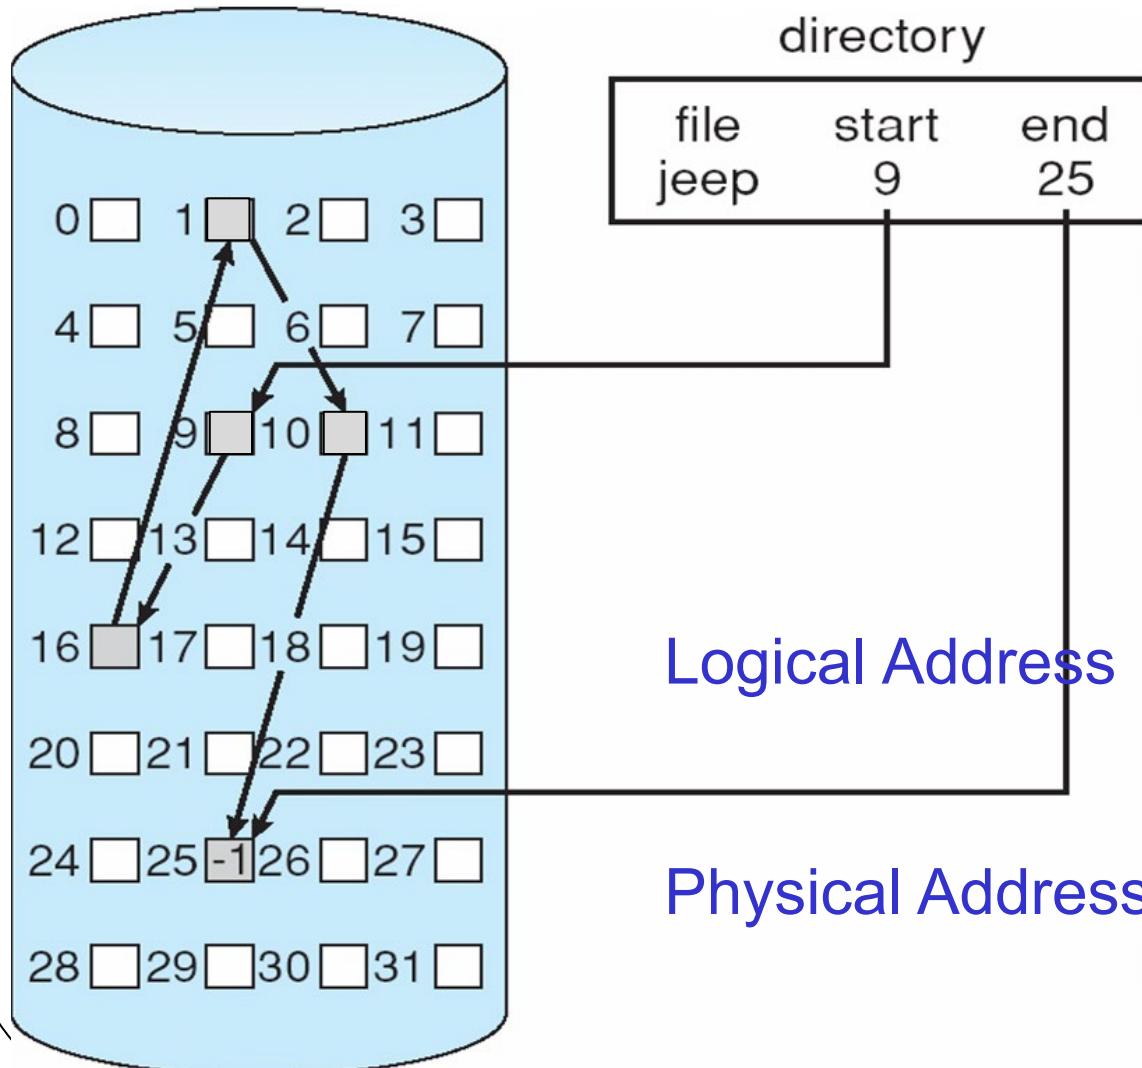
Linked Allocation

- Each file is a linked list of disk blocks; blocks may be scattered anywhere on the disk.
 - Simple: need only starting address
 - No waste of space
 - No constraint on file size: blocks can be allocated as needed
 - Problem: random access not supported

Linked Allocation (Cont.)

- Logical to Physical Address Mapping:
 - Suppose block size is 512 bytes, and first 4 bytes is reserved for the *pointer* to the next block in the list.
 - Logical Address $\Rightarrow Q \times 508 + R$
 - Block to be accessed is the $(Q+1)$ th block in the linked chain of blocks representing the file
 - Displacement into block = $R + 4$

- Allocate as needed, link together; e.g., file starts at block 9



Indexed Allocation

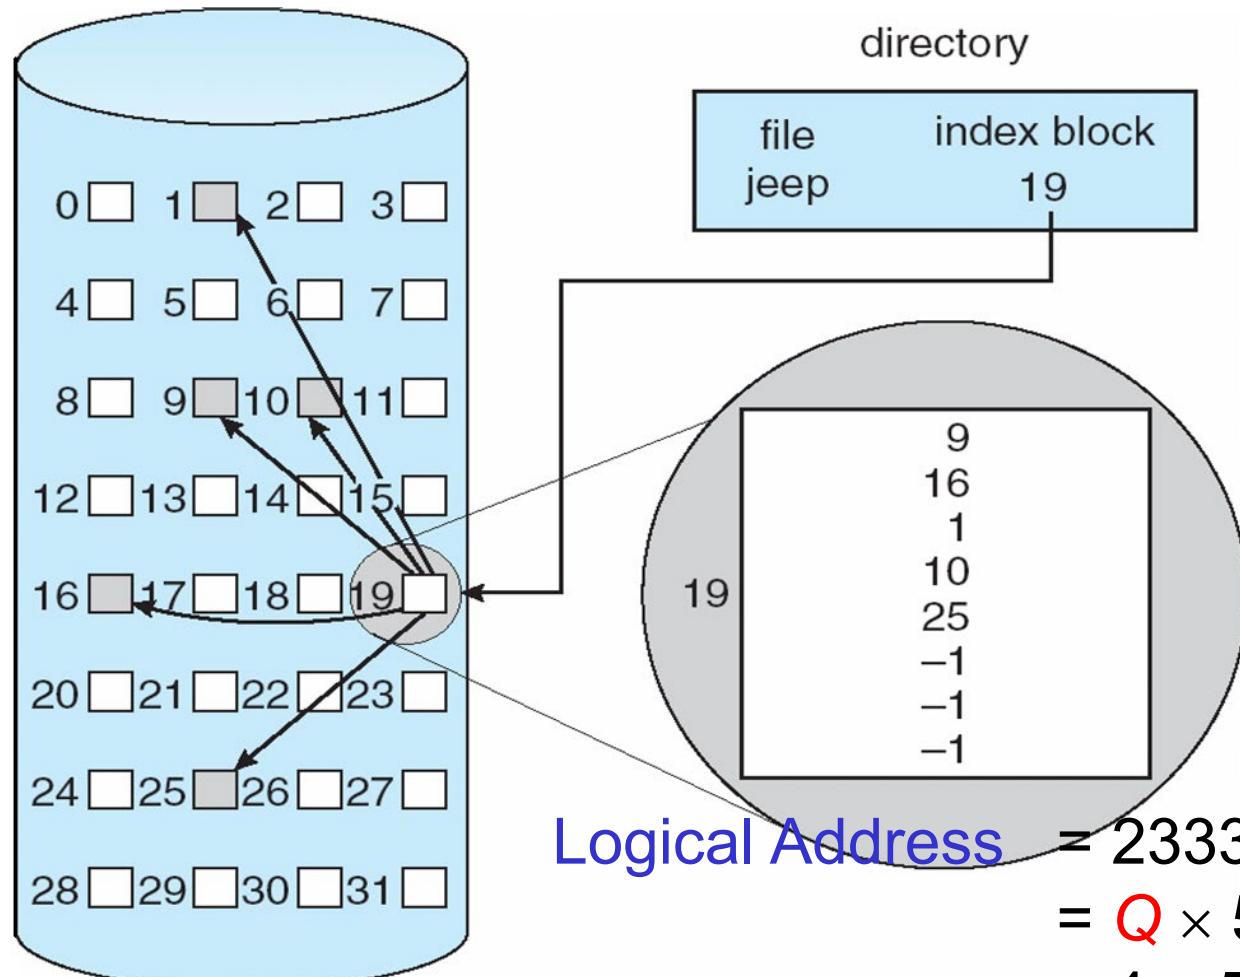
- Each file has an index block which contains all pointers to the allocated blocks. Directory entry contains the block number of the index block.
 - Support random access.
 - Dynamic storage allocation without external fragmentation (similar to the allocation of main memory using paging scheme).
 - *Problem:* Overhead of keeping index blocks and address mapping.

Indexed Allocation (Cont.)

- Logical to Physical Address Mapping
 - Suppose maximum size of a file is 128K bytes and block size is 512 bytes.
 - 2 blocks are needed for index table (4 bytes are used for each pointer)
 - Logical Address $\Rightarrow Q \times 512 + R$
 - Displacement into index table = Q
 - Displacement into block = R

Why are 2 blocks needed for index table in this case ?

Example of Indexed Allocation



Logical Address = 2333

$$= Q \times 512 + R$$

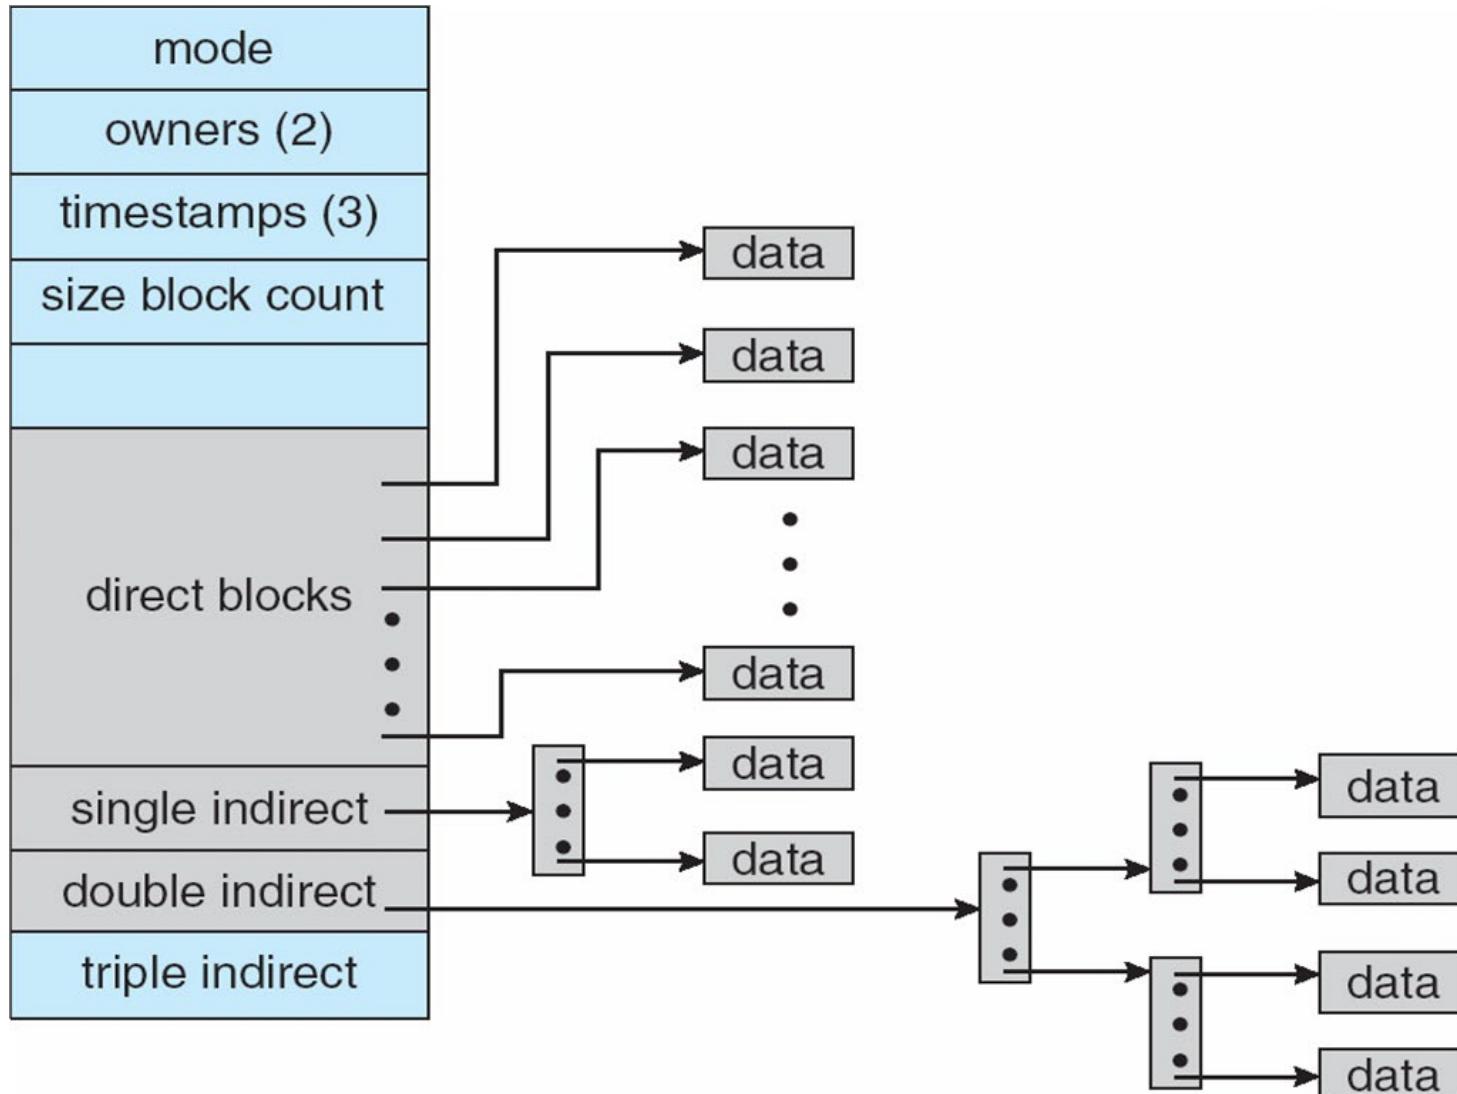
$$= 4 \times 512 + 285$$

Physical Address = 285th byte of block 25

Indexed Allocation: UNIX inode

- For each file or directory, there is an *inode* (index block).
- The inode contains:
 - *file attributes*
 - 12 pointers point to *direct blocks* (data blocks)
 - 3 pointers point to *indirect blocks* (index blocks). They are:
 - *single indirect*;
 - *double indirect*; and
 - *triple indirect*

inode (Cont.)



inode (Cont.)

- Assume 4-byte block pointer and 4K bytes block.
Maximum file size:
 - direct pointers: $12 \times 4K = 48K$ bytes
 - direct + single indirect pointers: $48K + 2^{22}$ bytes
 - Direct + single direct + double indirect pointers:
 $48K + 2^{22} + 2^{32}$ bytes > 4 gigabytes

$$2^{10} \times 2^{10} \times 4K = 2^{32}$$

$$\frac{2^{12}}{2^2} \times 4K = 2^{22}$$

of pointers per block

inode (Cont.)

- Each directory entry in UNIX contains an inode number and a file name.
- Each inode has a fixed location in disk.
- File look-up (i.e., search for an inode of a specific file) is straightforward. For example, looking up /usr/ast/mbox:

inode (Cont.)

Root directory

1	•
1	..
4	bin
7	dev
14	lib
9	etc
6	usr
8	tmp

Looking up
usr yields
i-node 6

I-node 6
is for /usr

mode
size
times

132

I-node 6
says that
/usr is in
block 132

Block 132
is /usr
directory

6	•
1	..
19	dick
30	erik
51	jim
26	ast
45	bal

/usr/ast
is i-node
26

I-node 26
is for
/usr/ast

mode
size
times

406

I-node 26
says that
/usr/ast is in
block 406

Block 406
is /usr/ast
directory

26	•
6	..
64	grants
92	books
60	mbox
81	minix
17	src

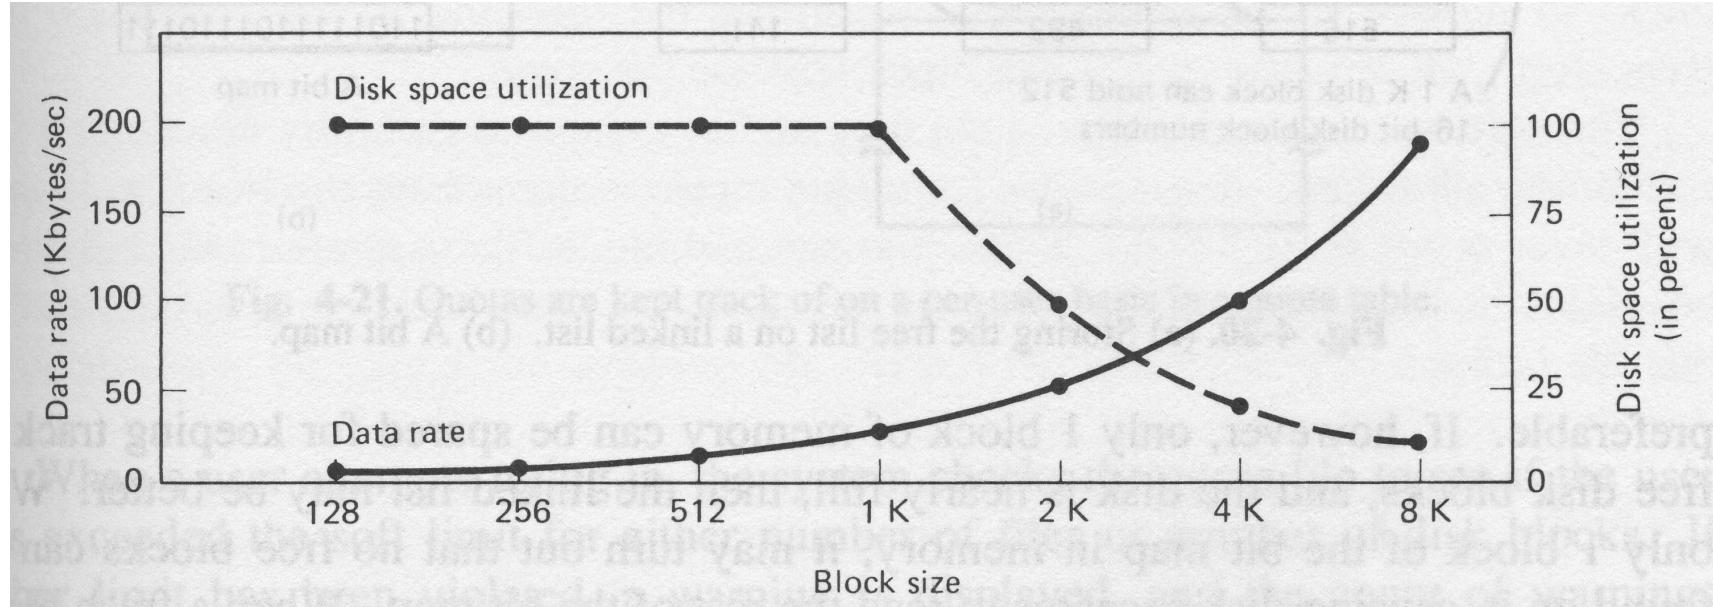
/usr/ast/mbox
is i-node
60

- Assume initially only root directory is in memory.

Disk-Space Management

- Block size affects both data rate and disk space utilisation
 - **Big block size**: file fits into few blocks \Rightarrow fast to find & transfer blocks, but wastes space if file does not occupy the entire last block
 - **Small block size**: file may consist of many blocks \Rightarrow slow data rate
 - Trade-off between time and space utilisation has to be compromised

Disk-Space Management (Cont.)



Disk-Space Management (Cont.)

- Keeping Track of Free Blocks: Similar to the issue of *Keeping Track of Memory Usage* in memory management under variable partition multiprogramming. Methods are:
 - Bit Map or Bit Vector
 - Linked List

Bit Map or Bit Vector

- Each block is represented by 1 bit:
0 \Rightarrow block is free; 1 \Rightarrow block is allocated
- Bit map requires extra space, e.g.,
 - block size = 2^9 bytes
 - disk size = 2^{34} bytes (16 gigabyte)
 - bit map size = $(2^{34}/2^9)/8 = 2^{22}$ bytes

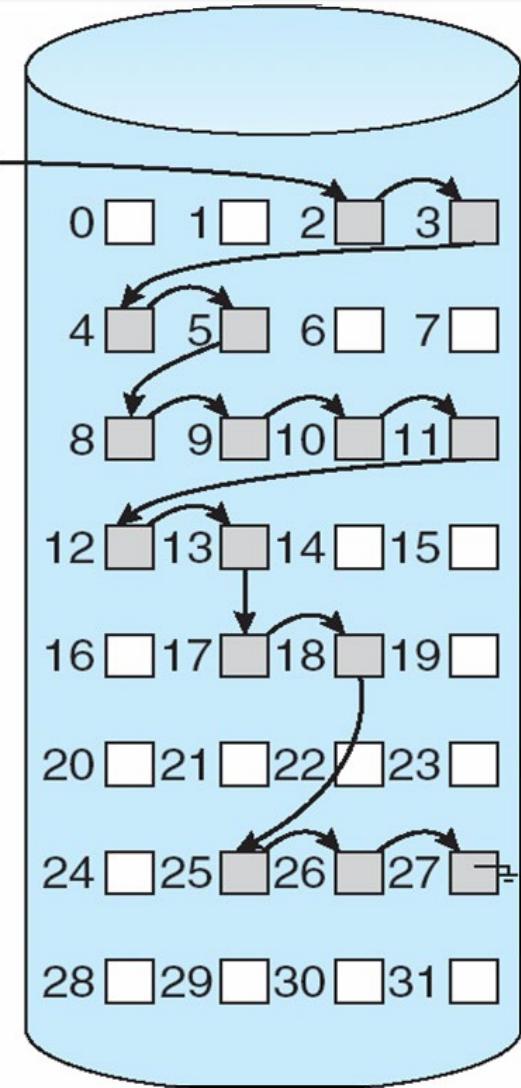
Bit Map (Cont.)

- Bit map is usually kept in a fixed place on disk. May be brought into memory for efficiency - needs to write back to disk occasionally for consistency and security.
- Easy to locate free blocks, but inefficiency unless the entire map is kept in memory.

Linked List

- Link together all the free disk blocks, keeping a pointer to the first free block in a special location on the disk and caching it in memory.
- No extra space required, but not efficient.

free-space list head



Summary

- Logical File System
 - Files
 - e.g., file structures,
 - Directories
 - e.g., support for file sharing
 - File Protection
- File Organization Module
 - Allocation Methods: contiguous, linked and indexed allocation
 - Disk Space Management